

Clear answers. Human judgment.

A focused prototype for families and the people running their center.

The customer problem and scope

Parents need a quick, dependable answer; operators need fewer repeat interruptions and a way to fix recurring gaps. This prototype finishes one loop: ask a question, inspect a cited policy, flag a gap, publish an improvement, and test again. Seven fictional policies cover everyday questions. Text and topic shortcuts keep the mobile experience simple; voice and document ingestion are intentionally out of scope.

Where AI helps, and where it stops

A small, quantized MiniLM model runs in a browser worker to match natural wording to staff-authored policy examples. Answers use the exact approved text, with a source and version. Semantic matches are labeled related guidance, not verified answers. Weak or ambiguous matches abstain. Topic-coverage checks prevent ordinary opening hours from answering holiday questions. Recognized emergencies and individual medical, custody, or account questions take a separate human-help path before retrieval. No model generates prices, grants exceptions, or authorizes pickup.

Operator control and visible quality

The staff inbox exposes knowledge gaps, related matches, requests for review, and negative feedback. Staff can draft, publish, edit, or withdraw policies and add examples in parents' words. New questions use the latest published version; earlier answers preserve their original citation. Counts distinguish explicit feedback from review status, and all three seeded questions are labeled samples. A staff-review action adds a local queue item; it does not pretend to send a real message.

Technical choices and limits

React and TypeScript, a static Render deployment configuration, and browser-local storage keep the demo inexpensive and isolated for each reviewer. No API key, per-question inference fee, or cloud transmission of question text is needed. First use downloads model/runtime assets; network hosts receive normal request metadata. A timeout or model error falls back to labeled basic search. Recognized sensitive details are not retained; basic email/phone redaction is incomplete. The public staff switch is simulated access, not authentication. There is no shared database, real notification, or conversational memory; this is not ready for real parent data.

Validation, tradeoffs, and next steps

Codex assisted implementation and debugging. Validation includes 63 unit cases, a production build, 15 real-model retrieval cases, and browser checks of local inference, mobile layout, persistence, and the question-to-policy loop. Evaluation caught snow-day/illness confusion and drove a retrieval fix. This small development set is not an independent accuracy estimate; heuristic safety and retrieval remain limited. Next: a held-out set of de-identified parent questions, an operator pilot, then authenticated roles, tenant-isolated persistence, policy approval history, real handoffs, and retention controls.

Source and run instructions: github.com/paulboutin/brightwheel-ai-front-desk

Context: [brightwheel family communication](#) - saving staff time while keeping family communication visible and organized.