

Clear answers. Human judgment.

Public policy answers for families. Personal questions belong in their secure portal.

Scope: a public front desk, not a parent portal

This prototype helps prospective and enrolled families ask general questions about a fictional center. The assignment leaves the access model open; a public front desk is a deliberate scope choice. Nine policies cover hours, tuition, meals, nut awareness, gluten-free requests, tours, emergencies, and medication administration. Child-specific conversations belong in the center's existing authenticated parent portal. That portal is outside this prototype: no login, child records, or secure teacher messaging is implemented or connected.

Where AI helps, and where it stops

A quantized MiniLM model runs in a browser worker to match natural wording to published policy examples. Answers use the exact approved text with a source and version. Semantic matches are labeled related guidance; weak or ambiguous matches abstain. General emergency and medication questions get policy answers. Reports of immediate danger receive urgent guidance even without a name. Personal care decisions point to the existing portal. Name and intent rules are incomplete; names are never proof of identity.

Parent experience and operator control

Parents open or close an in-page chat panel, switch topics without stacking conversations, and revisit saved questions. General unanswered questions can request a simulated email reply. Operators see gaps and feedback, reply, and turn a reviewed answer into an FAQ draft. Publication requires explicit review for all families; private replies never become knowledge automatically. New questions use current policies while earlier answers preserve their citations. Newly recognized personal or urgent questions retain only a routing label; no personal message is forwarded.

Technical choices and honest boundaries

React, TypeScript, static Render hosting, and browser-local storage make each reviewer's demo inexpensive and independent. The staff switch and email delivery are simulations, not protected accounts or real notifications. In a production portal, verified parent/child relationships and staff permissions would identify the sender and limit access; a secure staff screen alone cannot authenticate a public sender. No question text goes to a cloud model. First use downloads model/runtime assets; failures fall back to labeled basic search. Do not enter real family data; detection and redaction are incomplete.

Validation, tradeoffs, and next steps

Codex assisted implementation and debugging. Validation includes 143 unit cases, a production build, 30 real-model retrieval cases, and browser checks of mobile layout, persistence, publication, and routing. Parent testing exposed lunch-assistance overclassification and missing nut-free/gluten-free answers. Fixes add explicit dietary guidance and examples without guaranteeing an individual meal is safe; uncovered diets still abstain. This small development set is not an independent accuracy estimate. Next: a held-out evaluation and operator pilot, then an authorized existing-portal integration, durable tenant-isolated storage, audit history, delivery status, and retention controls. No portal URL is invented for this fictional center.

Source and run instructions: github.com/paulboutin/brightwheel-ai-front-desk

Context: [brightwheel family communication](#) - saving staff time while keeping family communication visible and organized.